

## Contents

|    |                                    |    |
|----|------------------------------------|----|
| 1  | Multi-Threading Basics             | 2  |
| 2  | Monitor Synchronization            | 3  |
| 3  | Specific Synchronization Mechanism | 4  |
| 4  | Threats of Concurrency             | 5  |
| 5  | Thread Pools                       | 7  |
| 6  | Async Programming                  | 8  |
| 7  | Memory Model                       | 10 |
| 8  | GPU Parallelization:               | 11 |
| 9  | Performance Optimizations          | 13 |
| 10 | Cluster Parallelization            | 14 |
| 11 | End                                | 14 |

# 1 Multi-Threading Basics

**async VS. parallel** *Attention:* Asynchron (Concurrent) ist not the same as parallel (see 1). Concurrent means that one core runs multiple threads. One thread makes progress while the other thread waits for IO. Parallel means that n cores run n threads at the exactly same time.



Figure 1: Async vs Parallel

**User-Level Threadings** User-Level Threads are implemented in a kind of runtime. The runtime itself runs in a process. Using this approach no real parallelism is possible, only one process).

**Kernel-Level Threading** Today you normally work with Kernel-Level Threads. The thread mechanism is implemented directly in the kernel. The programming language only interacts with the kernel API.

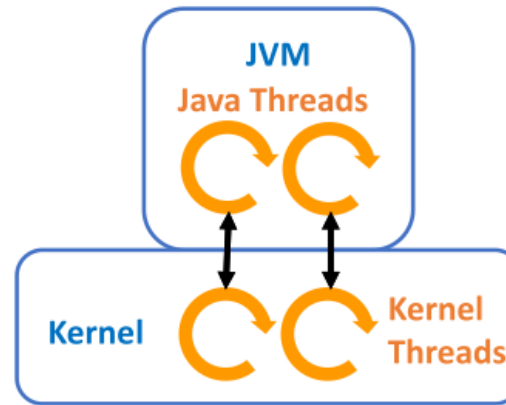


Figure 2: Kernel-Level Threads

**Processor Sharing** In modern system you have more threads than processors. That implies that not every thread can work at any time it wants. Therefore, the threads must *share* the processor.

**Multi-Tasking** The processor has two strategies how the processor sharing should work.

- cooperative
- preemptive

The preemptive is the used model on the processor. However, the most changes are cooperative because the thread has to wait for other resources (memory, IO, ...).

**Cooperative Multi-Tasking** In this model the thread works as long as it wants. If the thread want to wait, it must initiate the process by itself. A scheduler can not interrupt a running thread!

This model is not very common on modern systems.

**Preemptive Multi-Tasking** The scheduler uses a *Timer Interrupt* to interrupt a running thread. Each Thread can work for a specific max. interval. After that interval is over and the thread is not finished, the thread is interrupted and is added to the *Ready-Queue*. This is the common model which is today used.

## Thread State

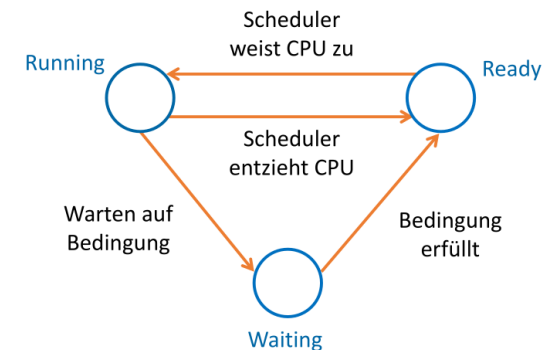


Figure 3: Thread States

**JVM Thread Model** The JVM is a single process system and has many threads (main, daemon threads, ...).

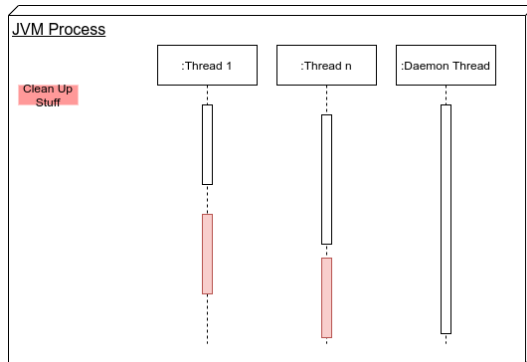


Figure 4: The JVM Thread Model

**JVM Termination** The JVM runs as long as at least one non-daemon thread is running. If the last non-daemon thread is terminated then the JVM is shutdown and all daemon threads are killed uncontrolled. Therefore, it is not possible to implement something like good-bye message when the thread terminates.

If inside a thread an uncaught exception occurs the other threads are **not** terminated and keep working.

**InterruptedException Java** In Java `join` and `sleep` can throw (theoretically) a checked `InterruptedException`. This exception is never thrown except you call `myThread.interrupt()`. Therefore, this exception should only be used if you implement a cooperative canceling. Otherwise, `xx.interrupt()` and the exception are indices for a code smell.

## 2 Monitor Synchronization

**synchronized keyword** In Java every object has a lock called Monitor-Lock. At the beginning of a **synchronized** block you acquire the lock as long the lock is free. If the lock is already acquired then you have to wait. After the **synchronized** block the lock is release (in **every** case)

```
public class Test {
    synchronized void do_something() {
        /* do stuff */
    }

    static synchronized void
        do_other_stuff() { }
}

public class Test2 {
    static void do_other_stuff() {
        synchronized (Test2.class) {}
    }

    void do_something() {
        synchronized (this) {}
    }
}
```

**Monitor lock** The monitor is a synchronization mechanism which use the *Wait & Signal* mechanism.

1. Fight for lock

- winner: enters the monitor (acquire the lock)
- others: wait for entrance

2. Thread in monitor does work

- if done: leaves monitor and wakes up other threads (**signal** / `signalAll()`)
- condition not satisfied: leaves monitor and goes to the right and wake up other threads (Wait on signal)

3. Go to 1.

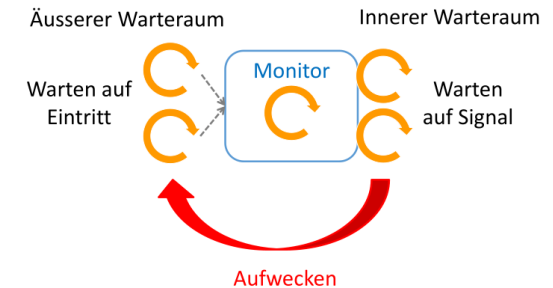


Figure 5: A monitor

**Wait Function in Loop** The function `wait()` must be always called inside a loop. In a simple if condition the following could happen:

1. thread 1: acquire lock, but condition is not fulfilled -> `wait()`
2. thread 2: acquire lock, change something, but thread 1 condition not fulfilled
3. thread 1: acquire lock again, make progress even the condition is not fulfilled

**notify vs notifyAll** A single `notify` can only be used when:

1. every thread has the same condition
2. only one **single** thread can make progress (One-In/One-Out)

**Thread Wake Up** A thread can be woken up by

1. notifyAll(), notify()
2. InterruptedException
3. Spurious Wake up (falsely wake up POSIX Thread API)

### Implement Monitor Lock

```
class BankAccount {
    private int balance = 0;

    public synchronized void deposit(int amount) {
        notifyAll();
        balance += amount;
    }

    public synchronized boolean withdraw(int amount) throws InterruptedException {
        int timeout = 0;
        while (amount > balance) {
            if (timeout >= 3) {
                return false;
            }
            wait();
        }

        balance -= amount;
        return true;
    }

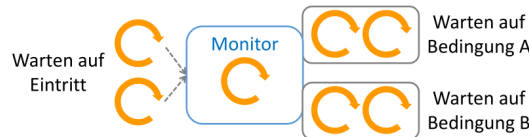
    public synchronized int getBalance() {
        return balance;
    }
}
```

## 3 Specific Synchronization Mechanism

**Semaphore** A semaphore is a synchronisation mechanism. A semaphore is basically a counter which can not count down below 0.

- **acquire():**
  - acquire resource from semaphore
  - wait, as long the counter is  $\leq 0$
  - otherwise, decrease counter
- **release():**
  - release resource
  - increment counter

**Lock and Condition** In lock & condition you have a monitor, but instead of one queue you have a queue for each condition. This has the benefit that not always all threads must be woken up. Only the threads where the condition may be now fulfilled.



:

fig:lock-and-conditions

### Lock and Condition Example

```
import java.util.concurrent.locks.
    Condition;
import java.util.concurrent.locks.Lock;
import java.util.concurrent.locks.
    ReentrantLock;
```

```
public class WarehouseWithLockCondition
{
    private Lock lock;
    private Condition nonEmpty;
    private Condition nonFull;

    public WarehouseWithLockCondition(
        int capacity, boolean fair) {
        this.capacity = capacity;
        store = 0;
        lock = new ReentrantLock(fair);
        nonEmpty = lock.newCondition();
        nonFull = lock.newCondition();
    }
}
```

```
@Override
public void put(int amount) throws
    InterruptedException {
    lock.lock();
    try {
        while (store + amount >
            capacity) {
            nonFull.await();
        }
        store += amount;
        nonEmpty.signalAll();
    } finally {
        lock.unlock();
    }
}
```

```
@Override
public void get(int amount) throws
    InterruptedException {
    lock.lock();
    try {
        while (store - amount < 0) {
            nonEmpty.await();
        }
        store -= amount;
        nonFull.signalAll();
    } finally {
        lock.unlock();
    }
}
```

```

    }
}

```

**Read-Write Locks** Mutual exclusion is too strong when only reading happens. Mutual exclusion is only required when minimal one thread wants to write.

|       | Parallel | Read | Write |
|-------|----------|------|-------|
| Read  |          | Yes  | No    |
| Write |          | No   | No    |

```

var rwLock = new ReentrantReadWriteLock(
    true);
rwLock.readLock().lock();
// read-only accesses
rwLock.readLock().unlock();
rwLock.writeLock().lock();
// write (and read) accesses
rwLock.writeLock().unlock();

```

**Problems With Read-Write Locks** The benefit of read-write locks is that you not have always mutual exclusion. But you can use the read lock only then, when **NO** write access happens on a shared object. But in common languages like Java it is not possible to encode the read only property into the API. However, Rust encodes the mutability of a function directly into the function arguments.

**Count-Down Latch** The count down latch is synchronisation mechanism. It's a counter which counts down to zero.

Using `await()` you wait until the counter is 0. Using `countDown()` you decrement the counter.

```

CountDownLatch waitForAll = new
    CountDownLatch(this.CARS);
protected void waitAllToBeReady()
    throws InterruptedException {
    waitForAll.await();
}

@Override
public void readyToStart() {
    waitForAll.countDown();
}

```

**Cyclic Barrier** A cyclic barrier is a synchronisation mechanism. The cyclic barrier is a meeting point for a fixed number of threads. You wait with the function `await()`. In contrast to the count down latch a cyclic barrier can be reused.

```

var gameRound = new CyclicBarrier(5);

/* 5 different players / threads */
while (true) {
    gameRound.await();
}

```

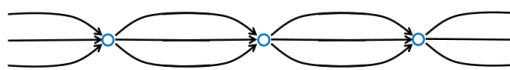


Figure 6: Cyclic Barrier

**Rendez-Vous** A rendez-vous is a special type of cyclic barrier with only 2 participants. Often you want to exchange data between those two threads. For this an **Exchanger** is used.

- Without exchange: `new CyclicBarrier(2)`
- With exchange: `Exchanger.exchange(something);`

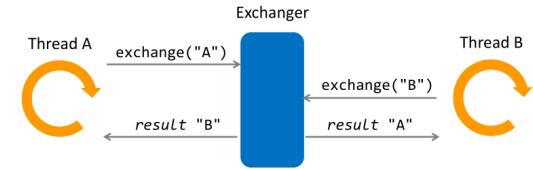


Figure 7: Exchanger Workflow

## 4 Threats of Concurrency

**Race Conditions** A race condition occurs then when multiple threads try to access the same resource without proper synchronization. This could result in wrong results or wrong behaviour. Often the reason is a data race. However, not always.

**Deadlock** A deadlock is when two or more threads waits on each other. This happens when thread 1 locks resource *a* and then resource *b* and thread 2 first resource *b* and then resource *a*.

```

/* Thread 1 */
synchronized(listA) {
    synchronized(listB) {
        listB.addAll(listA);
    }
}

/* Thread 2 */
synchronized(listB) {
    synchronized(listA) {
        listA.addAll(listB);
    }
}

```

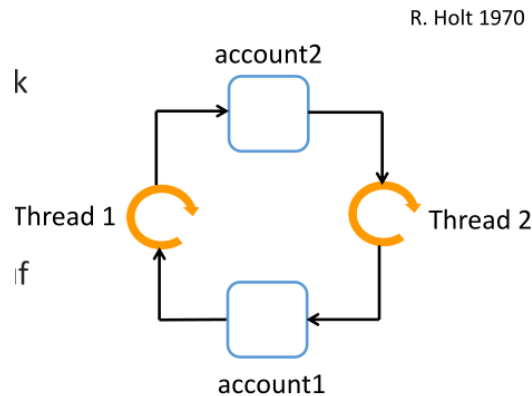


Figure 8: Deadlock detection using resource graph

**Starvation** When a thread never can acquire the lock for a resource, even when the lock is released from time to time, we call this *starvation*. Therefore, it is **not** a deadlock but a fairness problem.

**Data Race** When unsynchronized access happens (Race Condition) on the same memory location we call this a data race. For example, when a thread reads from a variable and another thread writes to the same variable. A read-read access is **not** data race (it is also not dangerous).

```
/* Thread 1 */
balance += 100;

/* Thread 2 */
balance += 50;
```

**Race Condition without data race** In the following code example we see, that no

data race is possible. However, a race condition is still possible. The read and the subsequent write are not atomic. The thread can be interrupted between the read and write action. And the action from another thread gets lost (Lost Update).

```
class BankAccount {
    int balance = 0;
    synchronized int getBalance() {
        return balance;
    }
    synchronized void setBalance(int x) {
        balance = x;
    }
}

/* On multiple threads */
account.setBalance(account.getBalance()
    + 100);
```

**No synchronization needed** Synchronization is not required when:

- the object is immutable (read only access)
- with confinement

**Confinement** Under confinement, we understand a structure that guarantees that only one thread can access an object at time.

- *Thread Confinement*: an object belongs to a single thread and is used only by this thread
- *Object Confinement*: an object is encapsulated in an already synchronized object

```
class ProductDatabase {
    private HashMap<String, Product>
        productMap = new HashMap<>();
```

```
public synchronized void addProduct(
    String name, String details) {
    productMap.put(name, new Product(
        details));
}

public synchronized
    getProductDetails(String name) {
    return productMap.get(name).
        getDetails(); // save,
        because String is immutable
}

public synchronized void notifySale(
    String name) {
    productMap.get(name).
        increaseSales();
}
}
```

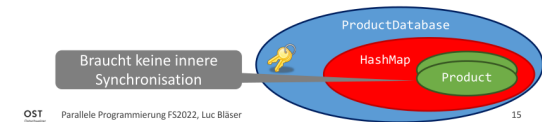


Figure 9: Example of object confinement

**Livelock** A live lock is a special case of a deadlock. While in a deadlock all relevant threads are sleeping in a live lock the threads are doing some waiting instruction. During these waiting instructions the CPU is working (busy wait).

**Avoiding deadlocks** To avoid deadlocks you have two possibilities:

- with a linear lock order
- coarse-grained locking (when linear lock order is not possible)

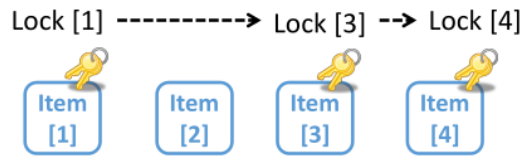


Figure 10: Linear Lock Order



Figure 11: Coarse Grained Lock

## 5 Thread Pools

**Thread Pool** A thread pool consists of a task queue and  $n$  worker threads. A new task is inserted in the queue. The new free thread from the thread pool takes the first task from the queue and process it.

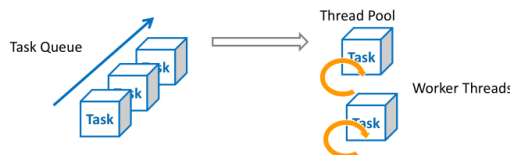


fig:thread-pool

**Number of Threads in Thread Pool** A thread pool has as much threads as processors (cores) and little bit more. A little bit more

than the number of processors is ideally because when one thread has to wait for I/O another thread can use the processor:

$$\begin{aligned} \# \text{Worker Threads} = \\ \# \text{Prozessoren} + \# \text{Pendente I/O-Aufrufe} \end{aligned} \quad (1)$$

**Thread Pool Java** Today in modern Java you should only use the `ForkJoinPool`. If an exception occurs then the exception is returned to the caller of `get()`.

```
var threadPool = new ForkJoinPool();
// submit is non-blocking
Future<Integer> future = threadPool.
    submit(() -> {
        int value = 1;
        // long calculation
        return value;
    });

// get blocks / same as await in Rust
Integer i = future.get();
```

**Task Parallelism .NET** In .NET you have one thread pool for:

- task parallelization
- data parallelization
- async programming

```
Task task1 = Task.Run(() => { /* Do some
    stuff */ });
task1.Wait(); // blocking

// Task with return value
Task task2 = Task.Run(() => { return
    3; });
Console.WriteLine(task2.Result); // blocking
```

```
// Task with Sub Tasks
Task.Run(() => {
    var left = Task.Run(() => Count(
        leftPart));
    var right = Task.Run(() => Count(
        rightPart));
    int result = left.Result + right.
        Result;
    return result;
});

// Parallele Statements
Parallel.Invoke(
    () => MergeSort(1, m),
    () => MergeSort(m, r)
);

// Parallel Loop
Parallel.ForEach(list, file => Convert(
    file));

// Parallel For – only if iterations are
// independent
Parallel.For(0, array.Length, i =>
    DoComputation(array[i]));

// PLINQ
from book in bookCollection.AsParallel()
where book.Title.Contains("
    Concurrency")
select book.ISBN;

from number in inputList.AsParallel().
    AsOrdered()
select IsPrime(number);
```

**Parallel Loop .NET** Often a loop has just a very short body. It is inefficient to execute every iteration in its own task. Therefore, the TPL groups automatically multiple bodies to a single task.

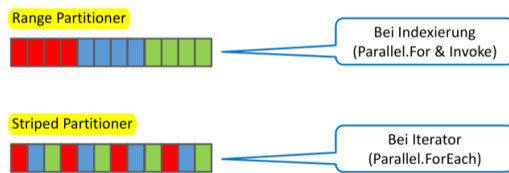


Figure 12: Partitioning in Parallel Loops

**Work Stealing Pool** Each Worker Thread has also its own task queue. A worker thread takes some task from the global queue and move it to the local queue. If the another worker thread wants to work but the global queue is empty, it steals it from other worker threads.

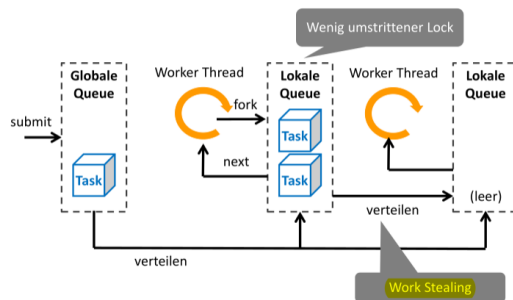


Figure 13: Work Stealing Thread Pool

**Thread Injection** When I injection new Threads in the thread pool during runtime is called *Thread Injection*. This can mitigate possible deadlocks when two task are not independent (only if not max threads is set).

**TODO Some Stuff**

## 6 Async Programming

**Types of asynchrony** You differ between two types of asynchrony:

- caller centric (pull)
  - The caller waits for task end and pulls the result
- callee centric (push)
  - task forwards the result to the successor task

### Continuation in C#

```
Task.Run(task1).
  ContinueWith(task2).
  COninueWith(task3);
```

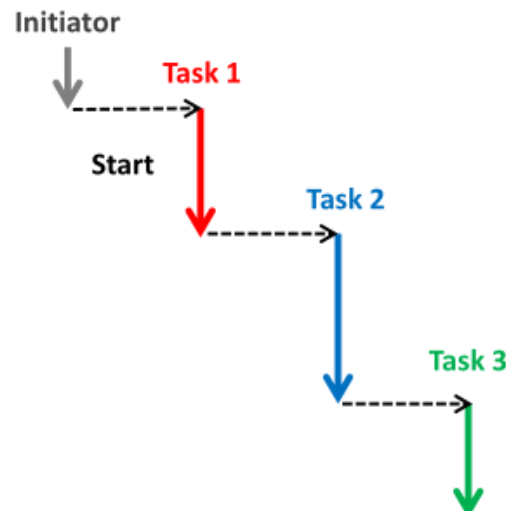


Figure 14: Task Continuations

```
Task.WhenAll(task1, task2).
  ContinueWith(continuation);
```

```
Task.WhenAny(task1, task2).
  ContinueWith(continuation);
```

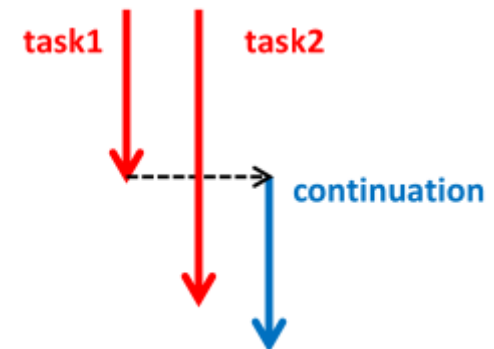
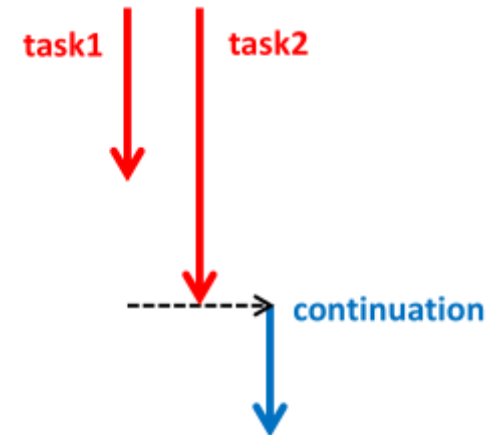


Figure 15: Multi Continuation



**Exception Handling Continuation** Exception in a Fire & Forget Tasks are ignored.

```
Task.Run(() => {
    throw new Exception("Some text"); //
    This exception is ignored by the
    system
});
```

To handle exceptions you could wait synchrony for finishing the task continuations (listening 1). Another option would be to subscribe for the `UnobservedTaskException`. However, this event is only triggered when the task is clean upped by the Garbage Collection. This can be at every point in time or even never (not deterministic).

```
task1.
    ContinueWith(task2).
    ContinueWith(task3).
    Wait();
```

Listing 1: Synchron Wait for finishing task continuation

## CompletableFuture Java

`CompletableFuture` is how you write async programs in modern Java and is the Java part of `Task` from C#.

The `CompletableFuture` has the methods listed in ??.

For Multi Continuation you have the following methods:

- `CompletableFuture.allOf(future1, future2)`
- `CompletableFuture.any(future1, future2)`

Other than .NET you can perform proper exception handling in continuations with `exceptionally()`.

**Update UI .NET** In WPF only the UI Thread is allowed to update the UI. Using a Dispatcher you can run an operation on the UI Thread from the background. In a WPF Project you can access the dispatcher:

- Code Behind: `this.Dispatcher`
- Other classes: `Application.Current.Dispatcher`

**Aysnc in .NET** A async function in .NET can only have the following return types:

- `Task`
- `Task<T>`
- `void` (fire and forget, should only be used in exceptional cases)

```
async Task<string> DownloadAsync(string url) {
    var web = new HttpClient();
    Task<string> task = web.GetStringAsync(url);
    string text = await task;
    return text;
}
```

} Synchron (Aufrufer Thread)  
} Asynchron (evtl. anderer Thread)

Figure 16: Async / Await Execution Modell

## async / await execution model .NET

The compiler splits the async function in multiple parts:

1. section before `await`: this section is executed synchrony
2. section after `await`: This is executed **after** the task is finished as a continuation.

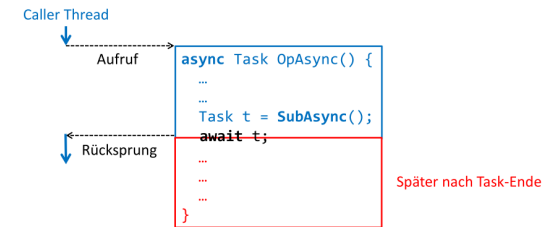
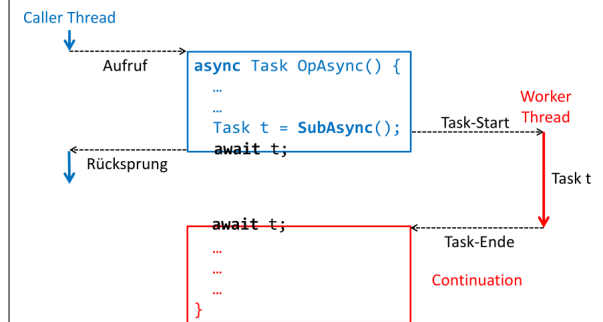
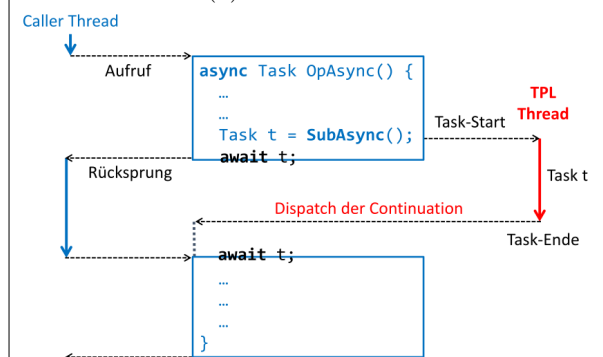


Figure 17: Async Method Call



(a) Non-UI Thread Execution



(b) UI Thread Execution

Figure 18: Thread Execution

## 7 Memory Model

**Lock Free Programming** Lock Free Programming is a type of programming where you write concurrent programs without any lock mechanism. In lock free programming you use the guaranties of the memory model. The goal of lock free programming is to create efficient synchronization.

**Weak Consistency** The memory access to the same memory address is seen differently by different threads. The possible output for j and i in 19 are:

- i = 0, j = 0
- i = 1, j = 1
- i = 0, j = 1
- i = 1, j = 0

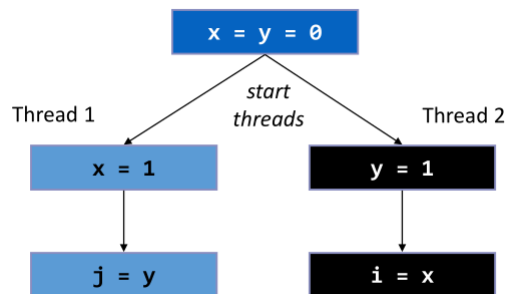


Figure 19: Example for Weak Consistency Example

The reason for this is that the compiler, runtime systems, CPUs, etc. can reorder instructions for optimization purpose. **Expect**, you use synchronization or memory barriers. This is called **weak consistency**.

**Atomicity Guaranty Java** Single reads / writes are atomic for:

- primitive data types until 32 bits
- object references
- long and double only with the **volatile** keyword

### Visibility Guaranty Java

Java guaranties the following visibility:

- changes before release are visible at acquire (20)
- changes until write are visible at read (21)
- the thread sees the correct start values and Join the output of the thread
- initialization of final variables (only relevant if you get the object from a data race!)

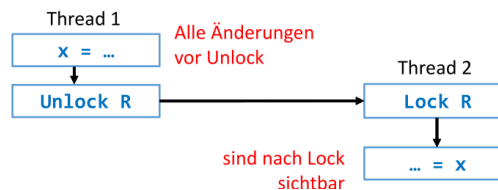


Figure 20: Visibility Lock / Unlock

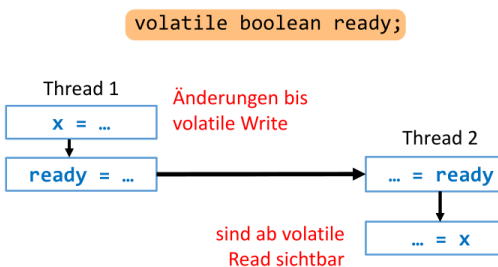


Figure 21: Visibility Volatile Write / Read

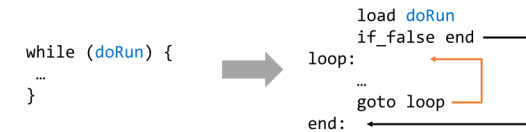


Figure 22: Visibility Problems

**Ordering Guaranty Java** The order of the visibility is the same as in visibility. Additional:

- synchronization instructions are never reordered to each other
- Lock/Unlock, volatile, thread start / join are never reordered
- if everything is a synchronization mechanism than we talk about total order

**volatile** boolean a = false, b = false;

Thread 1  
a = true;  
while (!b) { }

Thread 2  
b = true;  
while (!a) { }

Figure 23: Total Order Rendez-Vous

**Atomic Actions Java** Java has for each type an Atomic class. This class provides many functions which are atomic.

- getAndSet()
- compareAndSet()
- set()
- ...

**Optimistic Synchronization Java** In optimistic synchronization you perform the operation on a snapshot. If the snapshot is still the same value after the (long) operation you can update it. Otherwise, you have to do the same steps again until the fetch value is equal.

```
do {
    oldValue = var.get();
    newValue = calculateChanges(oldValue);
} while (!var.compareAndSet(oldValue,
    newValue));
```

```
// possible implementation for
// updateAndGet on the AtomicXXX
void updateAndGet(lambda
    calculateChanges) {
    do {
        oldValue = this.get();
        newValue = calculateChanges(
            oldValue);
    } while (!var.compareAndSet(oldValue,
        newValue));
}
```

**.NET Memory Model** In general the .NET Memory Model is the same as in Java with a few exceptions:

- atomicity: long / double are not atomic, event with volatile
- visibility: not defined, implicit defined by ordering
- ordering: only half and full fences

The .NET version of the Atomic classes in Java is the **Interlocked** class.

```
Thread.MemoryBarrier();
```

Listing 2: .NET Full Fence

```
volatile bool a = false, b = false;

Thread 1          Thread 2
a = true;         b = true;
Thread.MemoryBarrier();
while (!b) { }    Thread.MemoryBarrier();
                  while (!a) { }
```

Figure 24: Rendez-Vous in .NET

**Volatile Half Fences .NET** The **volatile** keyword in .NET is only a Half Fence. That means

- volatile writes: previous access stay before
- volatile reads: following access remain after

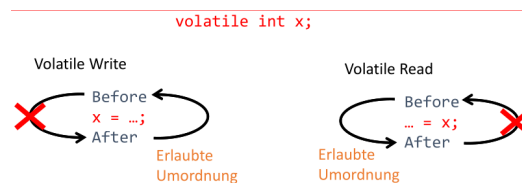


Figure 25: .NET Volatile Half Fences

## 8 GPU Parallelization:

### Latency vs throughput

**latency** how long does it take to execute a single instruction / operation

**throughput** number of instructions / operations completed per second

High throughput: lots of pizzas per hour

**compute & memory bound** The performance of a function is defined by memory bandwidth, compute bandwidth and latency. Total time for a function is **memory access + compute time**. If the compute time takes longer we say the function is *compute bound*. If the memory time takes longer we say the function is *memory bound*.

**When to use the GPU** The GPU is good if you can run the same instruction on multiple data at the same time. Therefore, GPU parallelism is only then a good idea when the function / code is compute bound. If the problem is the memory access then it is not a good idea to use the GPU.

**Arithmetic Intensity** The arithmetic intensity describes the ratio between computing and memory access. The higher, the better: efficient utilization of modern parallel processors.

$$\begin{aligned}
 t_c &> t_m \\
 \frac{ops}{BW_c} &> \frac{bytes}{BW_m} \\
 \frac{ops}{bytes} &> \frac{BW_c}{BW_m} \\
 \text{Arithmetic intensity} &= \frac{BW_c}{BW_m}
 \end{aligned} \tag{2}$$

If your arithmetic intensity is low, then you may not want to use the GPU. Instead, you want to optimize the memory access.

### Roof line model

The roof line model describes visually where your applications has performance issues.

In point  $O_1$  we can not fully use the GPU because the memory is too slow. However, in  $O_2$  the CPU is can not calculate faster and in this point it is a good idea to use GPU.

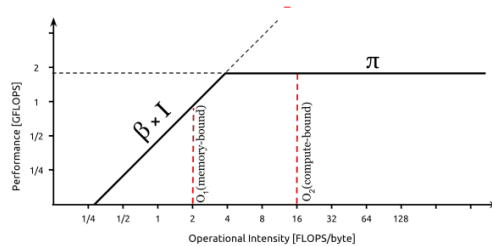


Figure 26: Roof line model

**GPU Structure** The big difference between the CPU and the GPU are the numbers of transistors for computing. The GPU has a lot more, but they are that good as the CPU ones.

A GPU consists of many **Streaming Multiprocessors (SM)**. And each SM consist of many **Streaming Processors (SP)** (Cores).

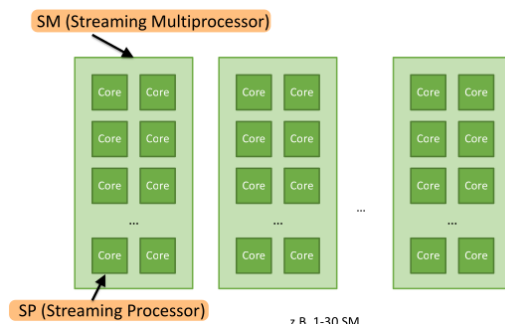


fig:gpu-structure

**SIMD** Single Instruction Multiple Data (SIMD) is performed on a Streaming Multiprocessor. Each core performs the exact same instructions but with different data.

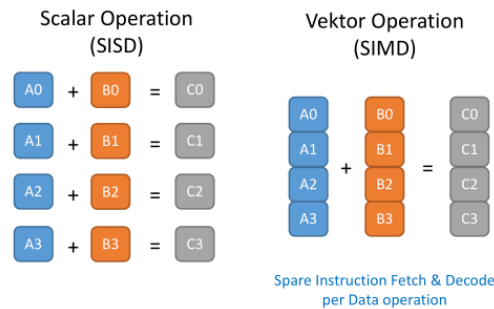


Figure 27: SIMD Vector parallelism

**NUMA** The NUMA model (Non-Uniform Memory Access) describes a model where two processors (CPU or GPU) do **NOT** share the memory. Therefore, one processor has to copy the data into other memory and back from there. The access time to the local memory (on-device) is much faster than the access to the remote one.

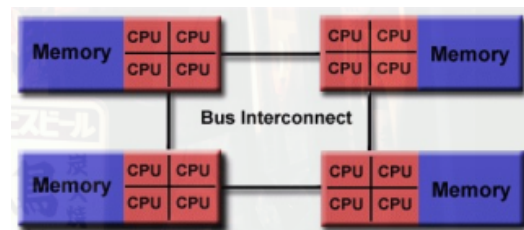


Figure 28: The NUMA model

### Thermilogy

| CUDA         | OpenCL       | Desc                              |
|--------------|--------------|-----------------------------------|
| Kernel       | Kernel       | Function to run on GPU            |
| Host Program | Host Program | Application which calls GPU       |
| Thread       | Work Item    | A single thread which runs on GPU |
| Block        | Work Group   | A collection of threads           |
| Grid         | NDRange      | A collection of blocks            |

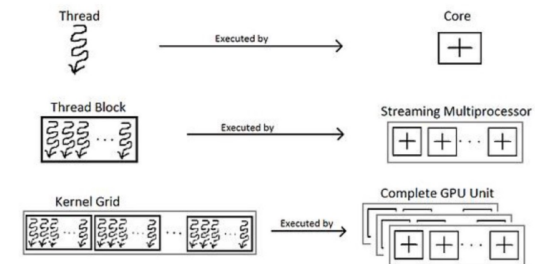


Figure 29: CUDA concepts on the GPU

**Thread Hierarchy** A Kernel is executed in a single thread. The single thread is executed simultaneously to other threads in a block. Many blocks are executed at the same time on a grid.

Internally, the threads are grouped into warps. These, are executed in a block.

Important: Each block can only execute a specific number of threads. And a grid can only execute a specific number of blocks.

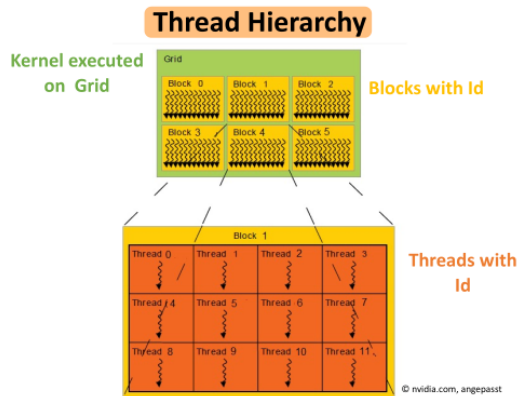


Figure 30: Thread Hierarchy on GPU

## 9 Performance Optimizations

**Compilation Workflow** The same applies to OpenCL (probably).

In CUDA you separate the device code from the host code. The device code is compiled into an assembly form (PTX code). The host codes loads and launches the compiled kernel (PTX code). The PTX code can be loaded by an application at runtime and is then **just-in-time compiled**. This has the benefits that you can benefit from the improvements of the compiler in the newer driver versions.

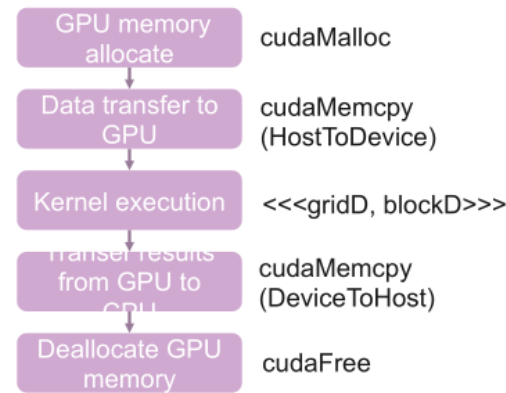


Figure 31: The CUDA execution

### Ceiling Integer

$$\text{blocks} = \frac{N + (m - 1)}{m} \quad (3)$$

$$\text{blocks} = \frac{N + 1023}{1024} \quad (4)$$

**Warp** A block is internally in **Warps** allocated. One Warp consists of 32 Threads. Each of these 32 threads executed the same instructions (same branch in same kernel).

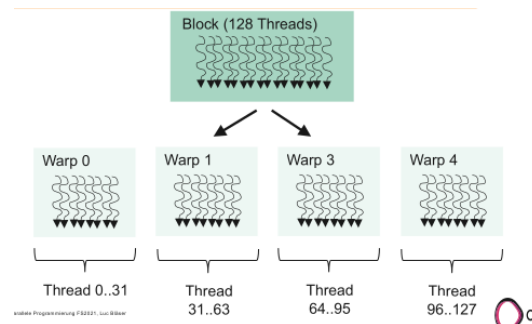


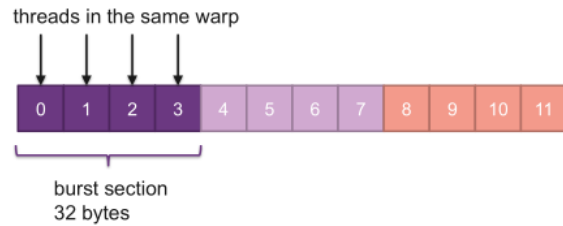
Figure 32: Block in Warp allocated

**Memory Coalescing** The pattern how a thread access the memory is critical. All threads in a warp should access the same are (burst). Then a combined access is possible. Otherwise, you have to do expensive individual accesses.

If you have a an access on memory use always the `threadIdx.x/y/z` to align the memory access.

```
data[(Expression without threadIdx.x) +
      threadIdx.x]
```

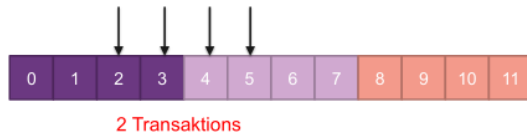
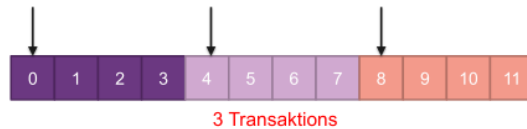
Listing 3: Align memory access (coalescing)



Read/write the burst in one transaction per warp  
burst section  
32 bytes

(a) Memory Coalescing

Misaligned => No Coalescing



(b) Non-Memory Coalescing

**CUDA Memory Model** All threads have access to the same global memory. Each thread block has shared memory, which is visible to all threads of the block. The shared memory is much faster because it is on-chip. The variables in the kernel are normally stored

in the register (fastest access). Each thread has private local memory which resides in device memory (slow). Only when all register are used the GPU uses the local memory.

Figure 2. Memory spaces on a CUDA device

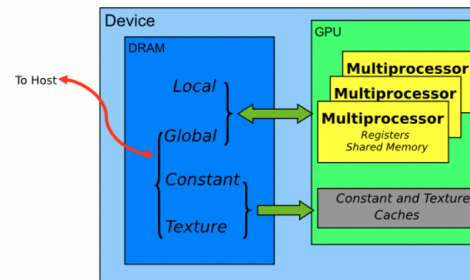


Figure 34: CUDA memory model

**Shared Memory** In matrix multiplication the same row / column is used multiple times. If the matrix are too big then can not load the whole row / column into memory. Therefore, the multiplication has to split up.

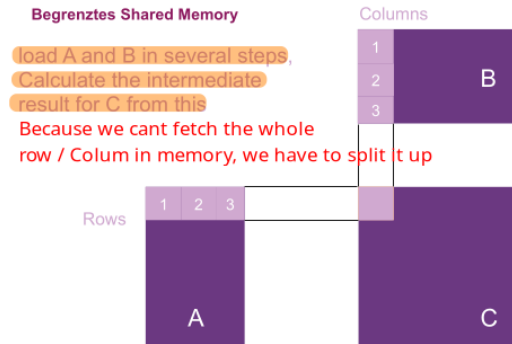


Figure 35: Tiling Matrix Multiplication

```
float sum = 0.0;
for (int tile = 0; tile < nofTiles; tile++) {
    // read tile from a and b in shared memory
    // each thread reads an element from each tile
    __syncthreads();
    // Multiply row of A-Tile by
    // Column of B-Tile from shared memory
    sum += partialProduct;
    __syncthreads();
}
C[row * M + col] = sum;
```

## 10 Cluster Parallelization

## 11 End